

DATA SOCIETY®

Advanced classification - day 2

*"One should look for what is and not what he thinks should be."
-Albert Einstein.*

Module completion checklist

Objective	Complete
Optimize random forest model	
Predict and evaluate the optimized model	
Optimize gradient boosting model	
Predict and evaluate the optimized boosting model	

Directory settings

- In order to maximize the efficiency of your workflow, you should encode your directory structure into variables
- Let the `main_dir` be the variable corresponding to your `af-werx` folder

```
# Set `main_dir` to the location of your `af-werx` folder (for Linux).  
main_dir = "/home/[username]/Desktop/af-werx"
```

```
# Set `main_dir` to the location of your `af-werx` folder (for Mac).  
main_dir = "/Users/[username]/Desktop/af-werx"
```

```
# Set `main_dir` to the location of your `af-werx` folder (for Windows).  
main_dir = "C:\\Users\\[username]\\Desktop\\af-werx"
```

```
# Make `data_dir` from the `main_dir` and  
# remainder of the path to data directory.  
data_dir = main_dir + "/data"
```

Loading packages

- Let's load the packages we will be using
- These packages are used for classification using decision trees, random forests, xgboost and other tools for classification best practices

```
import os
import pickle
import matplotlib.pyplot as plt
import numpy as np
import pandas as pd
from textwrap import wrap
from sklearn.model_selection import train_test_split
from sklearn import metrics
from sklearn import tree
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import accuracy_score
from sklearn.ensemble import GradientBoostingClassifier
from matplotlib.legend_handler import HandlerLine2D

# New today
from sklearn.model_selection import RandomizedSearchCV
```

Working directory

- Set working directory to `data_dir`

```
# Set working directory.  
os.chdir(data_dir)
```

```
# Check working directory.  
print(os.getcwd())
```

```
/home/[user-name]/Desktop/af-werx/data
```

Random forest vs gradient boosting

Goal for today

- Optimize a **random forest** and **gradient boosting** model to predict high yield in Costa Rican dataset
- Compare models at the end of today and decide on a **model champion**

Recap: random forest

- We understand how one **tree** works
- Why not try a **forest** ?

What is a random forest?

- Ensemble method used for **classification and regression tasks**
- Supervised learning algorithm which builds **multiple decision trees** and aggregates the result
- Uses a technique called Bootstrap Aggregation, commonly known as **Bagging**
- Limits overfitting and bias error

Random forest methodology

Step 1:
Set n & m
for forest



Step 2:
Build forest on
training data




Step 3:
Check
performance




Step 4:
Apply to test
data & evaluate



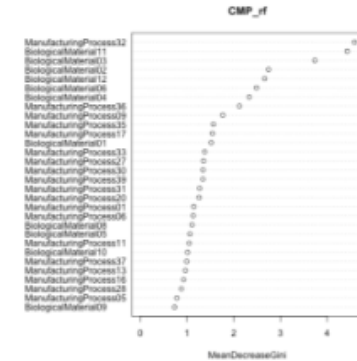

Step 5:
Use variable
importance as needed



	Act+	Act-
Pred+		
Pred-		



	Act+	Act-
Prod+		
Prod-		



Recap: why random forest?

- Reduction in overfitting
- Higher predictive accuracy
- Efficient with large datasets
- When should we use **decision trees** instead?
 - Intuitive and easily interpretable results
 - Less computationally expensive algorithm



Review data cleaning steps from last week

- **Today, we will be loading the cleaned dataset we used last class**
- To recap, the steps to get to this cleaned dataset were:
 - Remove household ID and individual ID
 - Remove variables with over 50% NAs
 - Transformed target variable to binary
 - Remove highly correlated variables

Load the cleaned dataset

- Let's load the dataset from last week, `costa_no_hc` - no highly correlated variables
- Save it as `costa_clean`

```
os.chdir(data_dir)
```

```
costa_clean = pickle.load(open("costa_no_hc.sav", "rb"))
```

```
print(costa_clean.head())
```

```
rooms  tablet  males_under_12  ...  urban_zone  age  Target
0      3      0              0  ...           1   43   False
1      4      1              0  ...           1   67   False
2      8      0              0  ...           1   92   False
3      5      1              0  ...           1   17   False
4      5      1              0  ...           1   37   False
```

```
[5 rows x 61 columns]
```

Print info on data

- Let's view the column names

```
costa_clean.columns
```

```
Index(['rooms', 'tablet', 'males under 12', 'males over 12',  
      'females under 12', 'females over 12', 'years of schooling',  
      'wall_block brick', 'wall_socket', 'wall_prefab cement', 'wall_wood',  
      'floor_mos_cer_terr', 'floor_wood', 'ceiling', 'electric_public',  
      'toilet_sewer', 'cookenenergy_elec', 'trash_truck', 'wall_bad',  
      'wall_reg', 'roof_bad', 'roof_reg', 'floor_bad', 'floor_reg',  
      'disabled_ppl', 'male', 'under10', 'free', 'married', 'separated',  
      'single', 'hh_head', 'hh_spouse', 'hh_child', 'num_65plus',  
      'dependency_rate', 'male_hh_head_educ', 'female_hh_head_educ',  
      'meaneduc', 'educ_primary_inc', 'educ_primary', 'educ_secondary_inc',  
      'educ_secondary', 'educ_undergrad', 'ppl_per_room', 'house_owned_full',  
      'house_owned_paying', 'house_rented', 'house_other', 'computer',  
      'television', 'num_mobilephones', 'region_central', 'region_Chorotega',  
      'region_pacifico', 'region_brunca', 'region_antlantica',  
      'region_huetar', 'urban_zone', 'age', 'Target'],  
      dtype='object')
```

Split into training and test sets

```
# Select the predictors and target.
X = costa_clean.drop(['Target'], axis = 1)
y = np.array(costa_clean['Target'])

# Set the seed to 1.
np.random.seed(1)

# Split into training and test sets.
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.3)
```

Tuning random forest model

- We saw when we built our random forest that there are many parameters
- Let's load our original model from last week
- Now we can take a look at the parameters currently in use

```
forest = pickle.load(open("model_forest.sav", "rb"))
```

```
forest.get_params()
```

```
{'bootstrap': True, 'class_weight': None, 'criterion': 'gini', 'max_depth': None, 'max_features':  
'auto', 'max_leaf_nodes': None, 'min_impurity_decrease': 0.0, 'min_impurity_split': None,  
'min_samples_leaf': 1, 'min_samples_split': 2, 'min_weight_fraction_leaf': 0.0, 'n_estimators': 100,  
'n_jobs': None, 'oob_score': False, 'random_state': 1, 'verbose': 0, 'warm_start': False}
```

Tuning random forest model

- The best approach would be to narrow the range of values for each parameter
- We can use `RandomizedSearchCV` to optimize our hyperparameters to sample from during fitting
- `GridSearchCV` can be computationally expensive, especially when dealing with a large hyperparameter space
- We can search only a random subset of parameter values instead, using `RandomSearchCV`

RandomizedSearchCV

- It uses a randomized search on hyperparameters
- Unlike GridSearchCV, it samples a fixed number of parameter settings from specified probability distributions

sklearn.model_selection.RandomizedSearchCV

```
class sklearn.model_selection. RandomizedSearchCV (estimator, param_distributions, n_iter=10, scoring=None, fit_params=None, n_jobs=None, iid='warn', refit=True, cv='warn', verbose=0, pre_dispatch='2*n_jobs', random_state=None, error_score='raise-deprecating', return_train_score='warn') ¶ \[source\]
```


Parameter grid

- Let's create a grid of parameter ranges

```
# Number of trees in random forest.
n_estimators = [int(x) for x in np.linspace(start = 200, stop = 2000, num = 10)]

# Number of features to consider at every split.
max_features = ['auto', 'sqrt']

# Maximum number of levels in tree.
max_depth = [int(x) for x in np.linspace(10, 110, num = 11)]
max_depth.append(None)

# Minimum number of samples required to split a node.
min_samples_split = [2, 5, 10]

# Minimum number of samples required at each leaf node.
min_samples_leaf = [1, 2, 4]

# Create the random grid.
random_grid = {'n_estimators': n_estimators,
               'max_features': max_features,
               'max_depth': max_depth,
               'min_samples_split': min_samples_split,
               'min_samples_leaf': min_samples_leaf}

print(random_grid)
```

Set up cross-validation function

- Now we instantiate the model, using 3-fold cross-validation with 100 different combinations

```
rf_random = RandomizedSearchCV(estimator = forest,
                               param_distributions = random_grid,
                               n_iter = 100,
                               cv = 3,
                               verbose = 0,
                               random_state = 1,
                               n_jobs = -1)

# Fit the random search model.
rf_random.fit(X_train, y_train)
```

```
rf_random.best_params_
```

```
{'n_estimators': 1000,
 'min_samples_split': 5,
 'min_samples_leaf': 2,
 'max_features': 'sqrt',
 'max_depth': 10}
```

Optimized random forest model

- Now we can use these optimized hyperparameters to implement the random forest again on `X_train`

```
optimized_forest = RandomForestClassifier(criterion = 'gini',  
                                         n_estimators = 1000,  
                                         min_samples_split = 5,  
                                         min_samples_leaf = 2,  
                                         max_features = 'sqrt',  
                                         max_depth = 10,  
                                         random_state = 1)
```

```
optimized_forest.fit(X_train, y_train)
```

```
RandomForestClassifier(bootstrap=True, class_weight=None, criterion='gini',  
                        max_depth=10, max_features='sqrt', max_leaf_nodes=None,  
                        min_impurity_decrease=0.0, min_impurity_split=None,  
                        min_samples_leaf=2, min_samples_split=5,  
                        min_weight_fraction_leaf=0.0, n_estimators=1000,  
                        n_jobs=None, oob_score=False, random_state=1, verbose=0,  
                        warm_start=False)
```


Knowledge Check 1



Exercise 1



Module completion checklist

Objective	Complete
Optimize random forest model	
Predict and evaluate the optimized model	
Optimize gradient boosting model	
Predict and evaluate the optimized boosting model	

Predict using the best model parameters

- Prediction class for each observation as a list

```
optimized_forest_y_predict = optimized_forest.predict(X_test)
# Look at the first few predictions.
print(optimized_forest_y_predict[0:5, ])
```

```
[False False False False False]
```

```
optimized_forest_accuracy = metrics.accuracy_score(y_test, optimized_forest_y_predict)
print ("Accuracy on test data (best model): ", optimized_forest_accuracy)
```

```
Accuracy on test data (best model): 0.8535564853556485
```

Optimized random forest: save final accuracy

- **Let's save our random forest score in our `model_final` dataset**
- We first have to load our `model_final` dataframe from last week

```
model_final_optimized =  
pickle.load(open("model_final_forest_gbm.sav", "rb"))
```


Optimized random forest: save final accuracy

```
# Add the model to our dataframe.
model_final_optimized = model_final_optimized.append({'metrics' : "accuracy" ,
                                                    'values' : round(optimized_forest_accuracy,
4),
                                                    'model': 'optimized forest' } ,
                                                    ignore_index = True)

print(model_final_optimized)
```

	metrics	values	model
0	accuracy	0.6046	knn_5
1	accuracy	0.6188	knn_GridSearchCV
2	accuracy	0.6287	knn_29
3	accuracy	0.6356	logistic
4	accuracy	0.7845	logistic_whole_dataset
5	accuracy	0.7859	logistic_tuned
6	accuracy	0.6611	tree_simple_subset
7	accuracy	0.9407	tree_all_variables
8	accuracy	0.7183	tree_all_variables_optimized
9	accuracy	0.9338	random forest
10	accuracy	0.8644	boosting
11	accuracy	0.8536	optimized forest

Knowledge Check 2



Exercise 2



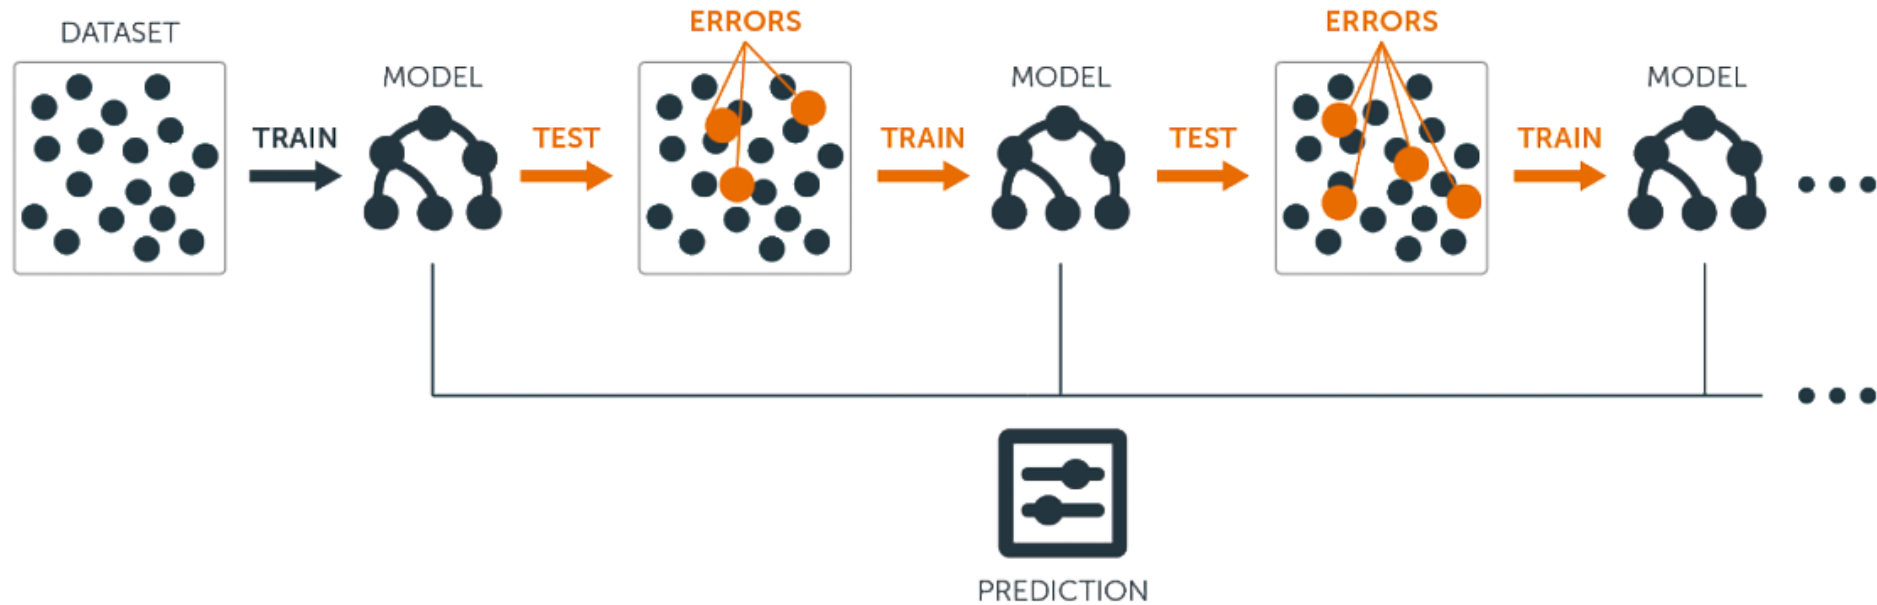
Module completion checklist

Objective	Complete
Optimize random forest model	✓
Predict and evaluate the optimized model	✓
Optimize gradient boosting model	
Predict and evaluate the optimized boosting model	

Recap: boosting intuition

- In simple linear regression, you can clearly see the residuals, which are the multiple points around the linear model
- Let's think of these residuals, but apply the concept to decision trees
- When **gradient boosting** uses decision trees, it follows these three steps:
 - it sees the **errors from a decision tree** on the dataset
 - it **identifies the pattern of the errors and builds a new decision tree** on them
 - it **repetitively leverages these patterns in residuals to strengthen the overall model**

Gradient boosted trees



Source

Recap: boosting process

- The process of **gradient boosting** is math heavy and complex
- However, for now, it can be simplified to three steps that we just discussed:
 1. Fit a decision tree model to the data
 2. Fit a decision tree model to the residuals
 3. Create a new model
- **Gradient boosting** can be used with classification or regression
- The generalization of the multiple weak learners occurs by the optimization of a differentiable **loss function**
- The **loss function** will change based on the model's target variable:
 - **Regression:** *gradient descent* used to minimize MSE
 - **Binary classification:** *logistic function*

Ways to optimize gradient boosting

- We saw when we built our boosting model that there are many parameters
- All the values of our original boosting model are set to the `GradientBoostingClassifier` defaults within `sklearn`
- We are now going to optimize the model **focusing on the four parameters we called out**
 - `max_depth = None`
 - `min_samples_split = 2`
 - `min_samples_leaf = 1`
 - `max_features = None`

Define an optimal number function

- Before we optimize individual parameters, let's build a function that will help us store the parameters we will be using in our `optimized_gbm`
- The inputs are:
 - `values` : list of values for given parameter that we iterate through
 - `test_results` : predictions on test set for each parameter that we iterate over
- The output is:
 - `best_value` : the actual parameter value that performs best and that we will use in our final optimized boosting model

```
# Define function that will determine the optimal number for each parameter.
def optimal_parameter(values, test_results):
    best_test_value = max(test_results)
    best_test_index = test_results.index(best_test_value)
    best_value = values[best_test_index]
    return best_value
```

Optimize: learning rate


```
# Learning Rate
learning_rates = [1, 0.5, 0.25, 0.1, 0.05, 0.01]
train_results = []
test_results = []

for eta in learning_rates:
    model = GradientBoostingClassifier(learning_rate=eta)
    model.fit(X_train, y_train)
    train_pred = model.predict(X_train)
    acc_train = accuracy_score(y_train, train_pred)
    train_results.append(acc_train)
    y_pred = model.predict(X_test)
    acc_test = accuracy_score(y_test, y_pred)
    test_results.append(acc_test)

optimal_learning_rate = optimal_parameter(learning_rates, test_results)
```

Plot: learning rate

```
from matplotlib.legend_handler import  
HandlerLine2D  
line1, = plt.plot(learning_rates,  
train_results, 'b', label= "Train accuracy")  
line2, = plt.plot(learning_rates, test_results,  
'r', label= "Test accuracy")  
plt.legend(handler_map={line1:  
HandlerLine2D(num_points=2) })  
plt.ylabel("Accuracy")  
plt.xlabel("Learning rate")  
plt.show()
```



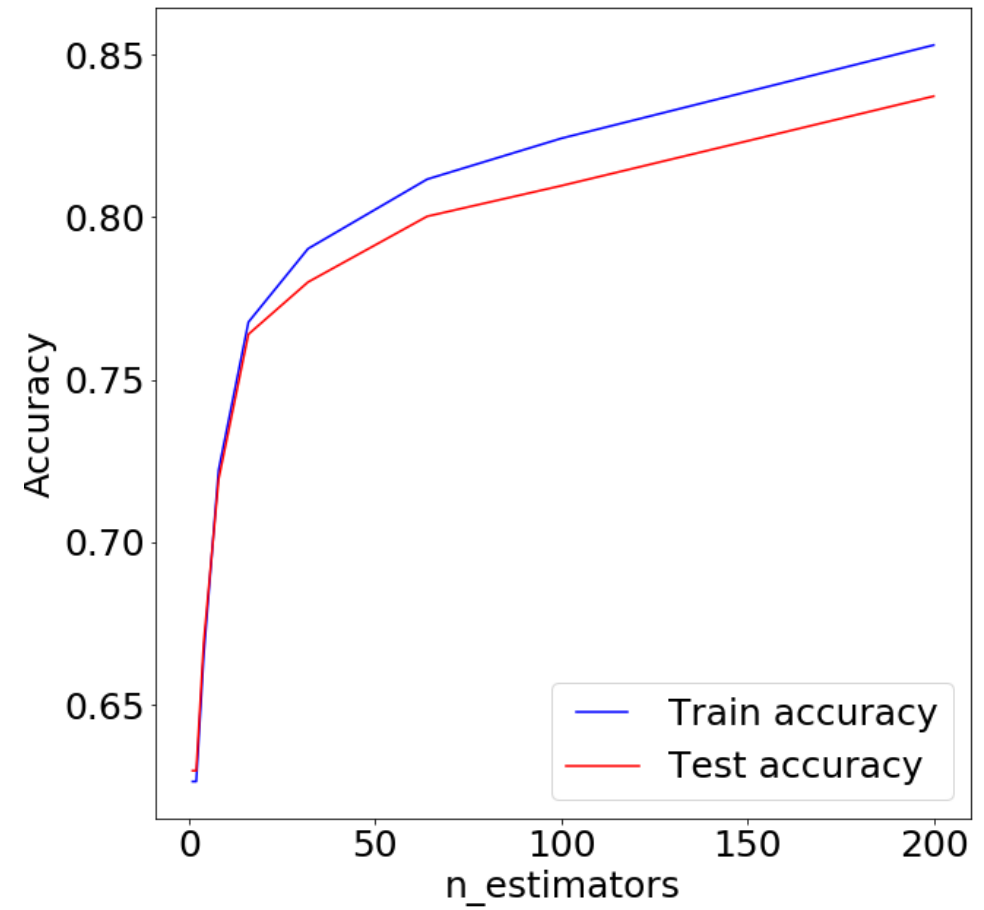
Optimize: n estimators

```
n_estimators = [1, 2, 4, 8, 16, 32, 64, 100, 200]
train_results = []
test_results = []
for estimator in n_estimators:
    model = GradientBoostingClassifier(n_estimators=estimator)
    model.fit(X_train, y_train)
    train_pred = model.predict(X_train)
    acc_train = accuracy_score(y_train, train_pred)
    train_results.append(acc_train)
    y_pred = model.predict(X_test)
    acc_test = accuracy_score(y_test, y_pred)
    test_results.append(acc_test)
```

```
optimal_n_estimators = optimal_parameter(n_estimators, test_results)
```

Plot: n estimators

```
from matplotlib.legend_handler import  
HandlerLine2D  
line1, = plt.plot(n_estimators, train_results,  
'b', label= "Train accuracy")  
line2, = plt.plot(n_estimators, test_results,  
'r', label= "Test accuracy")  
plt.legend(handler_map={line1:  
HandlerLine2D(numpoints=2) })  
plt.ylabel('Accuracy')  
plt.xlabel('n_estimators')  
plt.show()
```



Optimize: max depth

- max_depth indicates how deep the tree can be
- **The deeper the tree, the more splits it has, which captures more information about the data**
- **But remember, there is a fine line between a well fit model and an *overfit* model**
- In our original model, max_depth = None - now, we are going to fit a decision tree with depths ranging from 1 to 32 and plot the training and test accuracy


```
# Max depth:
max_depths = np.linspace(1, 32, 32, endpoint=True)
train_results = []
test_results = []
for max_depth in max_depths:
    model = GradientBoostingClassifier(max_depth=max_depth)
    model.fit(X_train, y_train)
    train_pred = model.predict(X_train)
    acc_train = accuracy_score(y_train, train_pred)
    train_results.append(acc_train)
    y_pred = model.predict(X_test)
    acc_test = accuracy_score(y_test, y_pred)
    test_results.append(acc_test)
```

```
# Store optimal max_depth.
optimal_max_depth = optimal_parameter(max_depths, test_results)
```

Plot: max depth

- Let's plot the max depth train_results and test_results
- This will allow us to see when the model starts overfitting on train as well as when the optimal test results are achieved
- **What observations can you make?**

```
# Plot max depth over 1-32.  
line1, = plt.plot(max_depths, train_results,  
                  'b', label= "Train accuracy")  
line2, = plt.plot(max_depths, test_results, 'r',  
                  label= "Test accuracy")  
plt.legend(handler_map={line1:  
    HandlerLine2D(num_points=2) })  
plt.ylabel('Accuracy')  
plt.xlabel('Tree depth')  
plt.show()
```



Optimize: min samples split

- `min_samples_split`: minimum number of samples required to split an internal node
 - This can vary between considering at least one sample at each node to considering all samples at each node
 - When we **increase this parameter, the tree becomes more constrained** as it has to consider more samples at each node
 - We will vary the parameter from 10% to 100% of the samples

```
min_samples_splits = np.linspace(0.1, 1.0, 10, endpoint = True)
train_results = []
test_results = []
for min_samples_split in min_samples_splits:
    model = GradientBoostingClassifier(min_samples_split = min_samples_split)
    model.fit(X_train, y_train)
    train_pred = model.predict(X_train)
    acc_train = accuracy_score(y_train, train_pred)
    train_results.append(acc_train)
    y_pred = model.predict(X_test)
    acc_test = accuracy_score(y_test, y_pred)
    test_results.append(acc_test)
```


```
# Store optimal min_samples_split.
optimal_min_samples_split = optimal_parameter(min_samples_splits, test_results)
```


Plot: min samples split

- Let's plot the min samples split
train_results and test_results

- **What observations can you make?**

```
line1, = plt.plot(min_samples_splits,  
train_results, 'b', label = "Train accuracy")  
line2, = plt.plot(min_samples_splits,  
test_results, 'r', label = "Test accuracy")  
plt.legend(handler_map = {line1:  
HandlerLine2D(num_points = 2)})  
plt.ylabel('Accuracy')  
plt.xlabel('min samples split')  
plt.show()
```



Optimize: min samples leaf

- `min_samples_leaf` is the minimum number of samples required to be at a leaf node
- This parameter is similar to `min_samples_split` except that **this parameter describes the minimum number of samples at the leafs - the base of the tree**

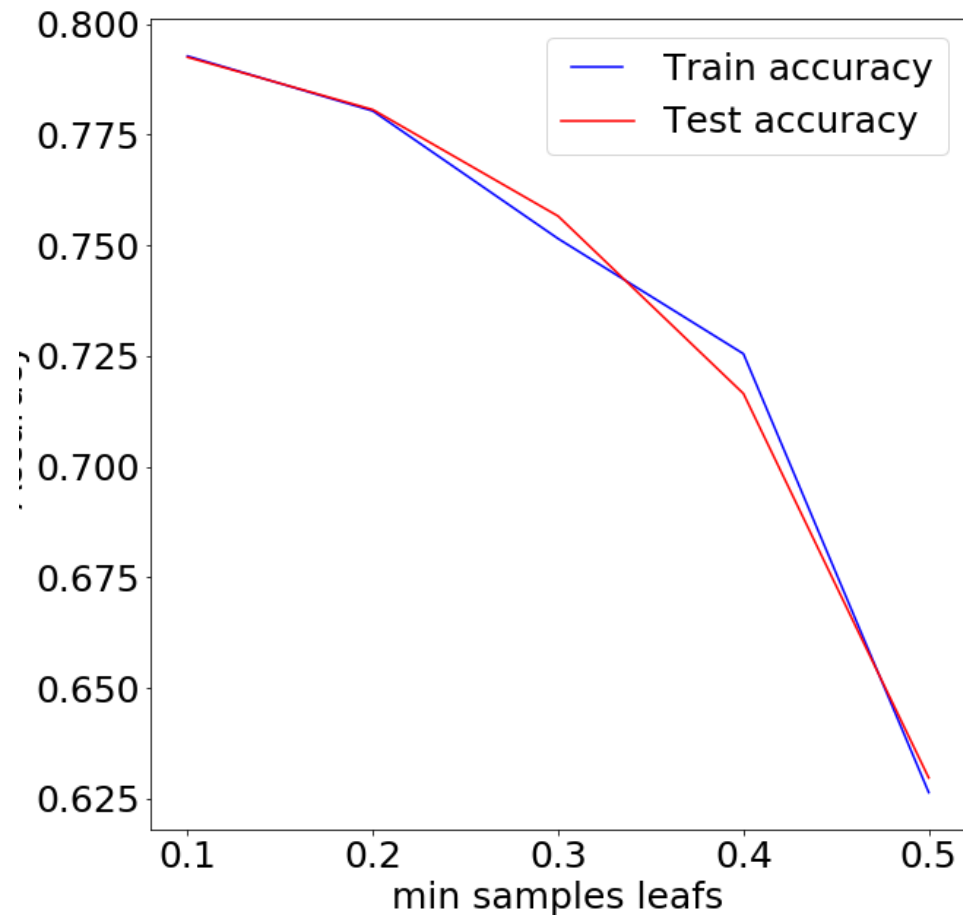
```
# Min_samples_leaf:
min_samples_leafs = np.linspace(0.1, 0.5, 5, endpoint = True)
train_results = []
test_results = []
for min_samples_leaf in min_samples_leafs:
    model = GradientBoostingClassifier(min_samples_leaf = min_samples_leaf)
    model.fit(X_train, y_train)
    train_pred = model.predict(X_train)
    acc_train = accuracy_score(y_train, train_pred)
    train_results.append(acc_train)
    y_pred = model.predict(X_test)
    acc_test = accuracy_score(y_test, y_pred)
    test_results.append(acc_test)
```

```
optimal_min_samples_leafs = optimal_parameter(min_samples_leafs, test_results)
```

Plot: min samples leaf

- Let's plot the min samples leaf
train_results and test_results
- **What observations can you make?**

```
line1, = plt.plot(min_samples_leafs,  
train_results, 'b', label = "Train accuracy")  
line2, = plt.plot(min_samples_leafs,  
test_results, 'r', label = "Test accuracy")  
plt.legend(handler_map = {line1:  
HandlerLine2D(num_points = 2)})  
plt.ylabel('Accuracy')  
plt.xlabel('min samples leafs')  
plt.show()
```



Optimize: max features

- `max_features` represents the number of features to consider when looking for the best split
- This parameter is set to `None` as its default value, so the tree will always look through all features
- This could sometimes cause overfitting and/or is computationally expensive when working with many variables

```
# Max_features:
max_features = list(range(1,X.shape[1]))
train_results = []
test_results = []
for max_feature in max_features:
    model = GradientBoostingClassifier(max_features = max_feature)
    model.fit(X_train, y_train)
    train_pred = model.predict(X_train)
    acc_train = accuracy_score(y_train, train_pred)
    # Add acc score to previous train results.
    train_results.append(acc_train)
    y_pred = model.predict(X_test)
    acc_test = accuracy_score(y_test, y_pred)
    # Add acc score to previous test results.
    test_results.append(acc_test)
```

```
optimal_max_features = optimal_parameter(max_features, test_results)
```


Plot: max features

- Let's plot the max features

train_results and test_results

- What observations can you make?

```
line1, = plt.plot(max_features, train_results,
                  'b', label = "Train accuracy")
line2, = plt.plot(max_features, test_results,
                  'r', label = "Test accuracy")
plt.legend(handler_map = {line1:
                          HandlerLine2D(num_points = 2)})
plt.ylabel('Accuracy')
plt.xlabel('max features')
plt.show()
```



Optimized model

- Now that we have now walked through four parameters that will help us optimize our gradient boosting model
- Let's look at what each of the optimal parameters are:

```
print("The optimal learning rate is:",  
      optimal_learning_rate)
```

The optimal learning rate is: 1

```
print("The optimal number of estimators is:",  
      optimal_n_estimators)
```

The optimal number of estimators is: 200

```
print("The optimal max depth is:",  
      optimal_max_depth)
```

The optimal max depth is: 3.0

```
print("The optimal min samples split is:",  
      optimal_min_samples_split)
```

The optimal min samples split is: 0.1

```
print("The optimal min samples leaf is:",  
      optimal_min_samples_leafs)
```

The optimal min samples leaf is: 0.1

```
print("The optimal max features is:",  
      optimal_max_features)
```

The optimal max features is: 59

Knowledge Check 3



Exercise 3



Module completion checklist

Objective	Complete
Optimize random forest model	✓
Predict and evaluate the optimized model	✓
Optimize gradient boosting model	✓
Predict and evaluate the optimized boosting model	

Build optimized model

- Now, we will run the optimized model on our X_train

```
# Set the seed.
np.random.seed(1)

# Implement the decision tree on X_train.
gbm_optimized = GradientBoostingClassifier(learning_rate = optimal_learning_rate,
                                           n_estimators = optimal_n_estimators,
                                           max_depth = optimal_max_depth,
                                           min_samples_split = optimal_min_samples_split,
                                           min_samples_leaf = optimal_min_samples_leafs,
                                           max_features = optimal_max_features)

# We can now see our optimized features where before they were just default:
print(gbm_optimized)
```

```
GradientBoostingClassifier(criterion='friedman_mse', init=None, learning_rate=1,
                           loss='deviance', max_depth=3.0, max_features=59,
                           max_leaf_nodes=None, min_impurity_decrease=0.0,
                           min_impurity_split=None, min_samples_leaf=0.1,
                           min_samples_split=0.1, min_weight_fraction_leaf=0.0,
                           n_estimators=200, n_iter_no_change=None,
                           presort='auto', random_state=None, subsample=1.0,
                           tol=0.0001, validation_fraction=0.1, verbose=0,
                           warm_start=False)
```

```
gbm_optimized_fit = gbm_optimized.fit(X_train, y_train)
```

Predict with optimized model

- Finally, let's predict on `X_test` and calculate our accuracy score
- **How is our optimized model doing?**
- **What other metrics can you also look at?**

```
# Predict on X_test.  
y_predict_gbm_optimized = gbm_optimized.predict(X_test)  
  
# Accuracy score.  
acc_score_gbm_optimized = accuracy_score(y_test, y_predict_gbm_optimized)  
  
print(acc_score_gbm_optimized)
```

```
0.8563458856345886
```

Add accuracy score to the final scores

- So we have it, let's add this score to the dataframe `model_final` that we created earlier
- Let's append the score to `model_final_optimized`

```
model_final_optimized = model_final_optimized.append({'metrics' : "accuracy" ,  
                                                    'values' : round(acc_score_gbm_optimized,4) ,  
                                                    'model': 'gbm_optimized'} ,  
                                                    ignore_index = True)  
  
print(model_final_optimized)
```

	metrics	values	model
0	accuracy	0.6046	knn_5
1	accuracy	0.6188	knn_GridSearchCV
2	accuracy	0.6287	knn_29
3	accuracy	0.6356	logistic
4	accuracy	0.7845	logistic_whole_dataset
5	accuracy	0.7859	logistic_tuned
6	accuracy	0.6611	tree_simple_subset
7	accuracy	0.9407	tree_all_variables
8	accuracy	0.7183	tree_all_variables_optimized
9	accuracy	0.9338	random forest
10	accuracy	0.8644	boosting
11	accuracy	0.8536	optimized forest
12	accuracy	0.8563	gbm_optimized

- Now that we have built all our models and have our final accuracy, let's discuss our results

Discuss model champion

- When we look for the best model, we do not just want to take the model with the highest accuracy score
- Remember to consider:
 - bias in models
 - optimized vs. biased models
- **We now want to open the class up for discussion about which model you would choose, and why**

Final scores

```
print(model_final_optimized)
```

	metrics	values	model
0	accuracy	0.6046	knn_5
1	accuracy	0.6188	knn_GridSearchCV
2	accuracy	0.6287	knn_29
3	accuracy	0.6356	logistic
4	accuracy	0.7845	logistic_whole_dataset
5	accuracy	0.7859	logistic_tuned
6	accuracy	0.6611	tree_simple_subset
7	accuracy	0.9407	tree_all_variables
8	accuracy	0.7183	tree_all_variables_optimized
9	accuracy	0.9338	random forest
10	accuracy	0.8644	boosting
11	accuracy	0.8536	optimized forest
12	accuracy	0.8563	gbm_optimized

```
pickle.dump(model_final_optimized, open("model_final_optimized_ensemble.sav", "wb" ))
```

Knowledge Check 4



Exercise 4



Module completion checklist

Objective	Complete
Optimize random forest model	✓
Predict and evaluate the optimized model	✓
Optimize gradient boosting model	✓
Predict and evaluate the optimized boosting model	✓

Workshop: next steps!

- Workshops are to be completed in the afternoon either with a dataset for a capstone project or with another dataset of your choosing
- Make sure to annotate and comment your code
- This is an exploratory exercise to get you comfortable with the content we discussed today
 - Step 1: Load data, clean data to get ready for exploratory data analysis
 - Step 2: Look for patterns by using visualizations and also unsupervised learning
 - Step 3: Using any insights gained from steps 1 and 2, move forward with building a base data table for modeling
 - Step 4: Determine your target variable, this will be the variable of interest, what you want to predict
 - Step 5: Build classification models using class materials as guidance, save accuracy scores like we did with `out_model_final` dataframes
 - Step 6: Determine a model champion

This completes our module
Congratulations!